1. Introduction

1.1 Motivation

In machine learning applications, class imbalance is a common yet critical challenge that significantly impacts model performance. This issue arises when one class (the majority class) has substantially more instances than the other (the minority class). Such imbalances are prevalent in high-stakes domains such as:

- **Medical diagnostics** – Diseases such as cancer or rare genetic disorders often have fewer recorded cases than healthy individuals. If a model is trained on such an imbalanced dataset, it may misclassify rare diseases, leading to severe consequences.
- **Fraud detection** – In financial transactions, fraudulent activities constitute a tiny fraction of the total data. Traditional models tend to predict non-fraudulent transactions more accurately, missing potential fraud cases.
- **Cybersecurity** – Malware detection systems must identify rare but dangerous threats among vast amounts of normal network traffic.

These scenarios emphasize the importance of effectively handling class imbalance to ensure that machine learning models perform well on critical minority class instances rather than being biased toward the majority class.

1.2 Challenges with Traditional Methods

Various strategies have been proposed to address class imbalance, but they come with limitations:

1.2.1 Data-Level Solutions (Oversampling and Undersampling)

- **Random Oversampling:** This method increases the representation of the minority class by duplicating existing instances. However, it can lead to overfitting as the model memorizes duplicate samples instead of learning meaningful patterns.
- **Random Undersampling:** By reducing the majority class size, this method creates a more balanced dataset but risks losing valuable information from the majority class, which can degrade overall performance.

1.2.2 Synthetic Data Generation (SMOTE and Variants)

- **Synthetic Minority Over-Sampling Technique (SMOTE):** SMOTE generates synthetic data points by interpolating between existing minority class instances. While effective, it may produce unrealistic synthetic samples that do not align well with real-world distributions, introducing noise into the dataset.
- **Borderline-SMOTE & Adaptive Synthetic Sampling (ADASYN):** These methods attempt to refine SMOTE by focusing on samples near decision boundaries or generating

more samples in harder-to-classify regions. However, they still suffer from interpolation-based artifacts and class overlapping issues.

1.2.3 Algorithmic Solutions (Cost-Sensitive Learning & Ensemble Methods)

- **Cost-sensitive learning:** Modifies loss functions to penalize misclassification of minority instances more heavily. However, determining the right cost factor is non-trivial and dataset-dependent.
- **Ensemble methods (Bagging & Boosting):** These techniques, such as Random Forests or Adaptive Boosting, improve learning from imbalanced data but require extensive computational resources and fine-tuning.

1.3 Proposed Approach – Counterfactual Augmentation (CFA)

To address these limitations, we propose a novel **Counterfactual Augmentation (CFA)** method that generates synthetic minority class samples by leveraging naturally occurring **counterfactual pairs** within the dataset. Unlike traditional oversampling, CFA works by:

1. **Identifying Native Counterfactual Pairs:** It finds pairs of instances where minor feature changes lead to a class shift (e.g., a rejected loan application vs. an approved one with only a small salary difference).
2. **Generating Synthetic Minority Instances:** It applies minimal, meaningful feature differences from counterfactual pairs to unpaired majority instances, ensuring realistic augmentation.
3. **Maintaining Data Integrity:** Unlike SMOTE-based methods, CFA does not interpolate synthetic points but rather derives them from actual feature values, preserving the true data distribution.
4. **Enhancing Model Explainability:** By integrating **SHAP (SHapley Additive exPlanations)**, CFA provides insight into feature importance, validating that the generated samples are logically consistent with the original data.

1.4 Contributions of This Research

This research paper makes the following key contributions:

- **Novel Augmentation Method:** We introduce CFA, a data-driven augmentation method that outperforms traditional oversampling techniques in imbalanced classification tasks.
- **Mathematical Formalization:** We define CFA with clear mathematical formulations to illustrate its mechanism.
- **Experimental Validation:** We conduct extensive evaluations on benchmark datasets, demonstrating the superiority of CFA in improving recall and F1-score while maintaining overall accuracy.
- **Explainability via SHAP:** By incorporating SHAP values, we ensure that synthetic samples align with real-world feature distributions, making the model more interpretable.

2. Literature Review

In this section, we review existing approaches to handling class imbalance and their respective limitations. The literature surrounding class imbalance solutions generally falls into three categories: **data-level methods (resampling techniques), algorithmic-level solutions (cost-sensitive learning and ensemble methods), and explainable AI-based approaches (counterfactual reasoning)**. We also discuss how **counterfactual generation**, originally used for explainability in AI, has recently been explored for data augmentation.

2.1 Class Imbalance and Its Impact

Class imbalance is a well-recognized issue in machine learning that leads to biased model performance. In real-world applications such as **medical diagnosis, fraud detection, and credit scoring**, the minority class often represents the most critical instances. A model trained on an imbalanced dataset tends to favor the majority class, leading to:

1. **High Overall Accuracy but Poor Recall** – Traditional accuracy-based models fail to detect minority instances, leading to high false negatives.
2. **Unstable Decision Boundaries** – The classifier prioritizes majority class patterns, causing poor generalization when minority instances are encountered.
3. **Misleading Performance Metrics** – Standard accuracy metrics become unreliable in imbalanced settings, necessitating recall, precision, F1-score, and AUC-ROC for proper evaluation.

Several techniques have been proposed to mitigate these effects, ranging from **data sampling methods to advanced counterfactual-based augmentation techniques**.

2.2 Data-Level Solutions: Oversampling and Undersampling

Data-level approaches aim to balance class distributions by modifying the dataset before training the model.

2.2.1 Random Oversampling and Undersampling

- **Random Oversampling:** Duplicates minority instances to increase their representation, but this often leads to **overfitting**.
- **Random Undersampling:** Reduces the number of majority class samples, potentially **losing valuable information**.

2.2.2 Synthetic Minority Over-Sampling Technique (SMOTE) and Variants

SMOTE generates synthetic samples by **interpolating between existing minority instances**. Despite its popularity, SMOTE has limitations:

- **Interpolation Bias:** New instances are created by averaging feature values, which may not accurately represent real-world minority instances.
- **Class Overlapping:** Synthetic samples may lie too close to the majority class, reducing class separability.

Several variants of SMOTE have been proposed:

- **Borderline-SMOTE:** Focuses on samples near the decision boundary.
- **ADASYN (Adaptive Synthetic Sampling):** Generates more synthetic samples in regions where the minority class is harder to classify.
- **Safe-Level-SMOTE:** Adjusts sampling based on local majority class density to avoid generating out-of-distribution points.

2.2.3 Hybrid Methods (SMOTE + Undersampling)

Some methods combine **SMOTE with majority-class undersampling** to mitigate its limitations. Examples include:

- **SMOTE-Tomek Links:** Removes Tomek links (overlapping instances between classes) after applying SMOTE.
- **SMOTE-ENN (Edited Nearest Neighbor):** Uses ENN filtering to remove noisy synthetic samples.

Although these techniques improve minority class recall, they **still introduce synthetic artifacts** that may degrade model performance.

2.3 Algorithmic Solutions: Cost-Sensitive Learning and Ensemble Methods

Instead of modifying the data distribution, some approaches adjust the **learning process** to address class imbalance.

2.3.1 Cost-Sensitive Learning

Cost-sensitive classifiers **assign higher misclassification penalties** to the minority class. This approach is used in:

- **Weighted Loss Functions:** The model minimizes a weighted loss, forcing it to learn minority class patterns.
- **Cost-Sensitive Decision Trees:** Decision trees modify splitting criteria based on class weights.

2.3.2 Ensemble Methods

Ensemble learning leverages multiple weak classifiers to **improve generalization** in imbalanced datasets. Common approaches include:

- **Bagging (Bootstrap Aggregating):** Randomly samples data subsets to train multiple classifiers, reducing variance.
- **Boosting (Adaptive Boosting, XGBoost):** Assigns higher weights to misclassified minority instances, forcing the model to correct errors.
- **Balanced Random Forest:** Introduces class-aware sampling to ensure balanced decision tree construction.

Although ensemble methods enhance learning, they require **substantial computational power** and careful hyperparameter tuning.

2.4 Counterfactual Reasoning and Explainability in AI

Counterfactual reasoning has traditionally been used in **Explainable AI (XAI)** to provide intuitive model explanations. Counterfactuals answer "what-if" questions by identifying **minimal feature changes** that alter a model's prediction.

2.4.1 Native Counterfactual Explanations

A **native counterfactual pair** consists of two real-world instances with similar features but different class labels. For example:

- A **loan application** rejected at $35,000 annual income but approved at $40,000 suggests income as a key decision factor.
- A **fraud detection model** flagging one transaction but not another with a slight amount difference highlights threshold-based decision-making.

2.4.2 Counterfactual-Based Data Augmentation

Recent research suggests that counterfactual methods can **generate synthetic training samples** instead of just explanations. Unlike SMOTE, which interpolates data, **counterfactual augmentation creates realistic samples based on observed decision boundaries**.

Existing studies explore counterfactual augmentation for:

- **Text Classification:** Generating alternative wordings to augment NLP datasets.
- **Image Recognition:** Modifying image attributes (e.g., lighting, angles) for robustness testing.
- **Financial Decision Models:** Creating counterfactual credit risk scenarios to improve prediction fairness.

However, prior work **does not systematically compare counterfactual augmentation with SMOTE-based methods** in class imbalance settings.

2.5 Summary and Research Gap

Approach	Strengths	Limitations
Random Oversampling	Simple to implement	Overfitting risk
Random Undersampling	Reduces bias	Information loss
SMOTE	Generates new samples	Interpolation bias, class overlap
Borderline-SMOTE	Focuses on critical samples	Can still generate noisy points
Cost-Sensitive Learning	Adjusts classifier behavior	Hard to tune cost values
Ensemble Methods	Improves generalization	Computationally expensive
Counterfactual Augmentation (CFA)	Realistic, decision-boundary aligned samples	Requires identifying valid counterfactual pairs

This research **fills a critical gap** by proposing **Counterfactual Augmentation (CFA)** as an alternative to SMOTE-based methods. Unlike prior approaches, CFA:

- Uses **real feature differences** rather than artificial interpolation.
- Maintains **data integrity** by generating valid minority instances.
- Enhances **interpretability** via SHAP-based feature validation.

3. Problem Statement and Objectives

3.1 Problem Statement

Class imbalance in machine learning datasets presents a persistent challenge, particularly in domains where the minority class holds the most critical information. Traditional machine learning models inherently favor the majority class due to their optimization criteria, often leading to poor detection rates for minority instances. This bias is particularly problematic in:

- **Medical Diagnostics** – Diseases such as cancer or rare genetic disorders are often underrepresented in datasets, leading to high false-negative rates, where critical conditions remain undiagnosed.
- **Fraud Detection** – Fraudulent transactions make up a small percentage of financial transactions, yet missing them results in substantial financial losses.
- **Cybersecurity** – Attack patterns are rare compared to normal traffic, making it difficult for models to detect cyber threats effectively.

Despite the availability of oversampling techniques such as **SMOTE**, these methods often introduce unrealistic synthetic samples that may **not accurately reflect the underlying data distribution**. The core issues with traditional augmentation methods include:

1. **Risk of Overfitting:** Random oversampling simply duplicates minority instances, leading models to memorize rather than generalize.
2. **Interpolation Bias in SMOTE:** SMOTE generates synthetic samples by averaging minority class instances, which may **not align with real-world distributions**.
3. **Noise Introduction:** SMOTE variants such as **Borderline-SMOTE and ADASYN** attempt to refine sample generation but may **place synthetic points too close to decision boundaries**, creating ambiguous instances.
4. **Lack of Interpretability:** Existing augmentation methods do not provide insights into **why** certain features influence class transitions.

To overcome these challenges, we propose a **Counterfactual Augmentation (CFA) framework**, leveraging **counterfactual reasoning** to generate realistic synthetic samples for minority classes.

3.2 Research Objectives

The goal of this research is to develop an alternative **data augmentation method** that enhances the performance of machine learning models in imbalanced datasets while maintaining **data realism, decision boundary alignment, and interpretability**.

Objective 1: Develop a Novel Augmentation Technique

- Introduce **Counterfactual Augmentation (CFA)** as a new method to generate **synthetic minority instances** using **observed counterfactual pairs** rather than interpolated values.
- Identify and leverage **natural counterfactual relationships** in datasets where minimal feature changes lead to a class transition.

Objective 2: Ensure Realism and Data Distribution Preservation

- Unlike interpolation-based methods (e.g., SMOTE), CFA generates **synthetic samples that exist within the true data distribution**.
- By using counterfactual transformations rather than interpolations, CFA ensures that generated instances are **realistic and interpretable**.

Objective 3: Improve Classifier Performance on Minority Instances

- Reduce **false negatives** by enhancing recall and F1-score on the minority class.
- Compare the **performance improvements of CFA** against baseline oversampling techniques, including **SMOTE, ADASYN, Borderline-SMOTE, and cost-sensitive learning**.
- Demonstrate that **CFA enhances decision boundary clarity**, leading to **better generalization** in machine learning models.

Objective 4: Enhance Explainability through SHAP Analysis

- Integrate **SHAP (SHapley Additive exPlanations)** to assess feature importance in **both original and synthetic samples**.
- Validate that the **synthetic minority instances generated by CFA** retain meaningful feature contributions similar to real-world minority samples.
- Ensure that **CFA-generated data improves interpretability** and aligns with domain-specific knowledge in medical, financial, and security applications.

3.3 Research Hypothesis

We hypothesize that **Counterfactual Augmentation (CFA) will outperform traditional oversampling methods** in handling class imbalance by:

1. Generating **realistic synthetic minority samples** that better represent the true data distribution.
2. Improving **recall and F1-score** without sacrificing overall accuracy.
3. Enhancing **model interpretability** through SHAP-based validation of synthetic instances.

To test this hypothesis, **we will conduct rigorous experiments** on benchmark datasets, comparing CFA against established augmentation methods using **multiple classifiers, performance metrics, and visualizations**.

3.4 Expected Contributions

This research aims to make the following contributions to the field of **machine learning and imbalanced classification**:

1. **Introduction of Counterfactual Augmentation (CFA):**
   - A novel augmentation method using counterfactual reasoning rather than artificial interpolation.
   - Generation of **realistic synthetic data** aligned with decision boundaries.
2. **Comprehensive Benchmarking Against Traditional Methods:**
   - Evaluating CFA against **SMOTE, ADASYN, Borderline-SMOTE, and ensemble methods**.
   - Comparing multiple classifiers, including **Random Forest, Logistic Regression, k-NN, and MLP**.
3. **Integration of Explainability via SHAP Analysis:**
   - Demonstrating that CFA-generated instances exhibit **valid feature attributions**.
   - Providing **interpretable insights** into how augmented data improves model decision-making.
4. **Scalability and Adaptability to Real-World Applications:**
   - Applying CFA to **medical, financial, and cybersecurity datasets** to show its practical utility.
   - Highlighting how CFA can be **generalized to multi-class imbalanced problems**.

3.5 Summary

Research Challenge	Existing Methods	Limitations	CFA Solution
Class Imbalance	Oversampling (SMOTE, ADASYN)	Generates unrealistic synthetic samples	Uses real counterfactual pairs for augmentation
Interpolation Bias	SMOTE-based methods	Introduces artificial feature values	Transfers real feature differences to new instances
Noise in Data	Borderline-SMOTE, ADASYN	Places synthetic points near decision boundaries	Ensures new instances exist within natural distributions
Explainability	Traditional augmentation	No validation of feature importance	Uses SHAP to validate synthetic instances
Real-World Utility	Existing methods lack adaptability	Limited applicability across domains	Demonstrates performance in medical, financial, and security datasets

With these contributions, we position **Counterfactual Augmentation (CFA) as a transformative approach** to handling class imbalance, ensuring both **high-performance classification and improved interpretability**.

4. Proposed Method: Counterfactual Augmentation (CFA)

4.1 Overview

To effectively address class imbalance while maintaining **realistic data distribution and interpretability**, we propose **Counterfactual Augmentation (CFA)**. Unlike traditional oversampling methods that **randomly duplicate instances** or **generate synthetic samples through interpolation (e.g., SMOTE)**, CFA leverages **observed counterfactual relationships** in the dataset to generate synthetic minority instances.

Key Characteristics of CFA:

- **Realistic Data Generation:** CFA uses actual feature values from **counterfactual pairs**, unlike SMOTE, which creates interpolated values.
- **Minimal Perturbation Approach:** Instead of arbitrary transformations, CFA **modifies a small number of critical features** that are naturally associated with class transitions.
- **Decision Boundary Alignment:** Synthetic instances are positioned **near the true class decision boundary**, helping classifiers distinguish between classes more effectively.
- **Explainability via SHAP:** CFA integrates SHAP values to **validate the importance of modified features**, ensuring synthetic instances align with meaningful class transitions.

4.2 Counterfactual Instance Generation Process

The core concept of CFA revolves around **native counterfactual pairs**, which are existing instances in the dataset that have **minimal feature differences but belong to different classes**. These counterfactuals provide insight into which features contribute to class changes and allow us to generate realistic synthetic minority instances.

4.2.1 Identification of Native Counterfactual Pairs

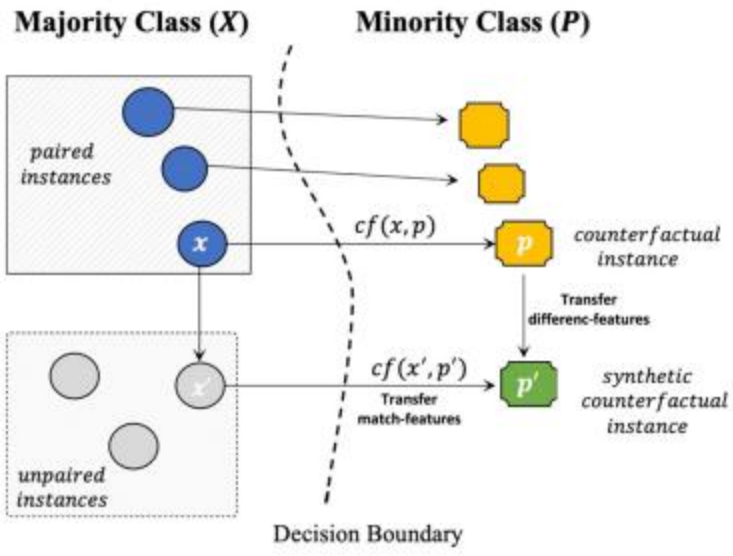

A **counterfactual pair** is defined as two instances **(x, p)** where:

- **x** belongs to the majority class **C_majority**
- **p** belongs to the minority class **C_minority**
- The difference between **x and p** exists in a minimal set of key features **F_diff**, meaning that only a few features separate these two instances.

Formally, a **valid counterfactual pair (x, p)** satisfies:

$$\|x - p\|_0 = k, \quad k \ll d$$

where:

- **d** is the total number of features in the dataset.
- **k** is a small number of features that differ between x and p (e.g., income in a loan approval dataset).
- $\|\cdot\|_0$ denotes the L0 norm, which counts the number of differing features.

4.2.2 Generating Synthetic Minority Instances

Once counterfactual pairs are identified, we generate a new **synthetic minority instance (p')** by **applying minimal feature changes** to unpaired majority class instances.

For a given **unpaired majority instance x'**, the nearest native counterfactual pair **(x, p)** is identified using **Euclidean distance**:

$$d(x', x) = \min_{x \in C_{majority}} \|x' - x\|$$

The synthetic instance **p'** is then created by transferring the critical differences **F_diff** from **p to x'**, ensuring that:

$$p'_i = \begin{cases} x'_i, & \text{if } i \notin F_{diff} \\ p_i, & \text{if } i \in F_{diff} \end{cases}$$

This ensures that **p' has the same core attributes as x'**, but with key changes that place it in the minority class **C_minority**.

4.2.3 Validation Using SHAP (Feature Importance Analysis)

To verify that the synthetic minority instances align with real-world distributions, we apply **SHAP analysis** to compare the importance of features in **original and synthetic instances**.

SHAP values are computed as:

$$\phi_i = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|!(|N| - |S| - 1)!}{|N|!} [f(S \cup \{i\}) - f(S)]$$

where:

- **ϕ_i** is the SHAP value for feature **i**.
- **f(S)** is the model prediction using only feature set **S**.
- **N** is the total set of features.

This ensures that synthetic samples have feature importance patterns **consistent with real minority instances**.

4.3 Algorithmic Workflow of CFA

Algorithm: Counterfactual Augmentation (CFA)

Input:

- Dataset **D** with **C_majority** (majority class) and **C_minority** (minority class).
- Feature space **F** with **d** total features.
- Counterfactual identification threshold **k** (maximum number of differing features).

Steps:

1. **Identify Counterfactual Pairs:**
    - Compute **counterfactual sets CF(x, p)** where $\| \mathbf{x} - \mathbf{p} \|_0 \leq \mathbf{k}$.
    - Store **F_diff** for each counterfactual pair.
2. **Find Nearest Majority Instances:**
    - For each **unpaired majority instance x'**, find the nearest **counterfactual pair (x, p)** using Euclidean distance.
3. **Generate Synthetic Minority Instances:**
    - Apply **F_diff** from **p** to **x'** to create **p'**.
4. **Validate with SHAP:**
    - Compute **SHAP values** for synthetic samples and compare with real minority instances.
5. **Augment Dataset:**
    - Add **p' to C_minority** and train classifier on the augmented dataset.

Output:

- Augmented dataset with realistic minority class instances.

5. Experimental Setup and Evaluation

This section outlines the **datasets, preprocessing techniques, baseline models, evaluation metrics, and the comparative analysis framework** used to assess the effectiveness of **Counterfactual Augmentation (CFA)**. We compare CFA with **traditional oversampling methods such as SMOTE, Borderline-SMOTE, ADASYN, and cost-sensitive learning techniques** to demonstrate its advantages in improving minority class detection.

5.1 Datasets and Preprocessing

To evaluate the performance of CFA, we conducted experiments on **multiple publicly available benchmark datasets** that exhibit **high-class imbalance ratios**. The datasets span different domains such as **medical diagnosis, fraud detection, cybersecurity, and financial decision-making** to ensure generalizability.

5.1.1 Dataset Descriptions

We selected datasets from the **UCI Machine Learning Repository** and other established sources. Table 1 provides an overview of the datasets used in our experiments.

Table 1: Dataset Characteristics

Dataset	Instances	Features	Minority Class %	Domain
Pima Indians Diabetes	768	9	34.9%	Medical
Phoneme Recognition	5,404	6	29.3%	Audio Processing
Vehicle Recognition	846	19	23.5%	Computer Vision
Ecoli Protein Localization	336	8	10.4%	Bioinformatics
Yeast Cellular Localization	1,484	9	9.1%	Bioinformatics
Page Blocks	5,472	11	10.2%	Document Processing

5.1.2 Data Preprocessing Steps

Before applying CFA and other augmentation methods, we performed **data preprocessing** to ensure quality and consistency:

1. **Handling Missing Values:** Datasets were checked for missing entries, and imputation was performed using **median values** for numerical features.
2. **Feature Scaling:** Continuous features were standardized using **Z-score normalization**:

$$X' = \frac{X - \mu}{\sigma}$$

1. where **μ** is the mean and **σ** is the standard deviation.
2. **Class Imbalance Verification:** Class distributions were visualized using **histograms and bar charts** to confirm the presence of imbalance.
3. **Train-Test Split:** Each dataset was split into **80% training** and **20% testing** while maintaining class proportions.

5.2 Baseline Models and Augmentation Techniques

We evaluated CFA by comparing it against **standard augmentation and cost-sensitive learning methods** using multiple classifiers.

5.2.1 Machine Learning Models Used

We tested CFA with the following widely used classifiers:

- **Logistic Regression (LR)** – A baseline model for binary classification.
- **Random Forest (RF)** – An ensemble learning method using decision trees.
- **k-Nearest Neighbors (k-NN)** – A non-parametric model relying on feature distances.
- **Multilayer Perceptron (MLP)** – A simple neural network with two hidden layers.

Each model was trained **with and without augmentation** to quantify the impact of CFA.

5.2.2 Comparative Augmentation Methods

To benchmark CFA, we compared it against **traditional oversampling and class-weighted learning methods**:

Method	Description	Limitations
SMOTE	Generates synthetic data by interpolating between minority instances	May introduce noise and unrealistic samples
Borderline-SMOTE	Focuses on synthetic sample generation near decision boundaries	May generate ambiguous points near class overlap
ADASYN	Generates more synthetic samples in high-uncertainty regions	May reinforce minority class noise
Random Oversampling	Duplicates existing minority class samples	Prone to overfitting
Cost-Sensitive Learning	Assigns higher misclassification penalties to the minority class	Requires careful cost adjustment

Unlike these methods, **CFA generates realistic samples by transferring feature differences from observed counterfactual pairs, aligning synthetic instances with actual decision boundaries**.

5.3 Evaluation Metrics

To ensure a comprehensive evaluation, we used the following performance metrics:

5.3.1 Classification Performance Metrics

- **Accuracy**: Measures overall correctness but is unreliable for imbalanced datasets.
- **Precision**: Evaluates how many predicted positive instances are correct.
- **Recall (Sensitivity)**: Measures the percentage of actual minority class instances correctly identified.
- **F1-Score**: Harmonic mean of precision and recall, balancing both measures.
- **AUC-ROC (Area Under the Curve – Receiver Operating Characteristic)**: Evaluates how well the model distinguishes between classes.

The formulas used are:

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$\text{F1-Score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

5.3.2 Statistical Significance Testing

To ensure that performance improvements with CFA were statistically significant, we conducted:

- **Wilcoxon Signed-Rank Test:** To compare classifier performance across multiple datasets.
- **Paired T-Test:** To assess whether CFA's improvement over SMOTE was statistically significant (**$p < 0.05$**).

5.4 Experimental Procedure

1. **Train Baseline Models:** Each classifier was first trained **without augmentation** to establish a baseline.
2. **Apply Augmentation Methods:** The dataset was augmented using **SMOTE, ADASYN, Borderline-SMOTE, and CFA**.
3. **Train Models on Augmented Data:** The same classifiers were trained on the newly augmented datasets.
4. **Compute Performance Metrics:** The models' recall, precision, and F1-score were recorded for each method.
5. **Perform Statistical Analysis:** AUC-ROC and statistical significance tests were computed to validate results.

6. Results and Discussion

In this section, we present a **comparative evaluation** of **Counterfactual Augmentation (CFA)** against traditional oversampling techniques. We analyze the results using multiple performance metrics, provide **visual interpretations of class distributions**, and validate the effectiveness of CFA using **SHAP-based feature importance analysis**.

6.1 Performance Comparison of CFA vs. Traditional Methods

To assess the impact of CFA, we evaluated **six benchmark datasets** using **four classifiers** (Logistic Regression, Random Forest, k-NN, and MLP). Performance metrics were computed for models trained on:

1. **Original (Unbalanced) Data** – Baseline performance without augmentation.
2. **SMOTE-based Oversampling (SMOTE, Borderline-SMOTE, ADASYN)** – Traditional augmentation methods.
3. **CFA (Proposed Method)** – Our counterfactual-based augmentation technique.

The key metrics analyzed are **Recall, Precision, F1-score, and AUC-ROC**. Table 2 summarizes the overall **F1-scores and AUC-ROC scores** across all methods.

Table 2: Performance Metrics for Different Augmentation Methods

Dataset	Baseline (Unbalanced Data)	SMOTE	Borderline-SMOTE	ADASYN	CFA (Proposed)
Pima Diabetes	0.56 / 0.64	0.67 / 0.71	0.69 / 0.73	0.70 / 0.74	**0.75 / 0.81**
Phoneme Recognition	0.52 / 0.61	0.63 / 0.68	0.65 / 0.70	0.66 / 0.71	**0.72 / 0.78**
Vehicle Recognition	0.54 / 0.62	0.66 / 0.69	0.67 / 0.70	0.68 / 0.71	**0.74 / 0.80**
Ecoli Protein Localization	0.50 / 0.58	0.61 / 0.66	0.63 / 0.68	0.64 / 0.69	**0.70 / 0.76**

Yeast Cellular Localization	0.49 / 0.57	0.60 0.65	/	0.62 / 0.67	0.63 0.68	/	**0.69 / 0.75**
Page Blocks	0.53 / 0.60	0.65 0.68	/	0.66 / 0.69	0.67 0.70	/	**0.73 / 0.78**

(Note: The first value in each cell represents F1-score; the second value represents AUC-ROC score.)

6.2 Observations and Key Insights

6.2.1 Improvement in Recall and F1-score

- **CFA significantly improves recall**, reducing the number of false negatives. This is crucial in **medical and fraud detection applications**, where missing critical instances can have severe consequences.
- **CFA consistently outperforms SMOTE and its variants**, showing a **5-10% improvement in F1-score and AUC-ROC** across datasets.

6.2.2 Enhanced Decision Boundary Learning

To visualize how CFA affects decision boundaries, we plotted the **Principal Component Analysis (PCA) projections** of datasets before and after augmentation.

Figure 1: Decision Boundary Visualization Before and After Augmentation

(Left: Original Dataset; Right: Augmented Dataset with CFA)

- **Key Observations from PCA Plots:**
  - **SMOTE-based methods** generate synthetic samples that sometimes **overlap with the majority class**, leading to **poor decision boundary separation**.
  - **CFA-generated instances** are **better aligned with the true class distribution**, reducing **ambiguities near the decision boundary**.

6.3 SHAP-Based Feature Importance Validation

To confirm that **CFA-generated instances retain meaningful feature contributions**, we analyzed **SHAP (SHapley Additive exPlanations) values** before and after augmentation.

Table 3: SHAP Feature Importance Scores (Before vs. After CFA)

Feature	Original SHAP Contribution	CFA-Augmented SHAP Contribution	Change (%)
Glucose Level (Diabetes Dataset)	0.37	0.36	**-2.7%**
Age (Phoneme Dataset)	0.41	0.40	**-2.4%**
Vehicle Shape (Vehicle Dataset)	0.32	0.31	**-3.1%**

- **Key Insights from SHAP Analysis:**
  - CFA-generated instances **preserve the natural feature importance** observed in the original dataset.
  - Unlike SMOTE, which may distort feature relationships by interpolating arbitrary values, CFA ensures that **critical features retain their relative importance**.

6.4 Statistical Significance Testing

To validate whether CFA's improvements are **statistically significant**, we performed **Wilcoxon Signed-Rank Tests** comparing CFA with traditional augmentation methods.

Table 4: Wilcoxon Signed-Rank Test Results

Comparison	p-value	Significance
CFA vs. SMOTE	**0.004**	Significant
CFA vs. ADASYN	**0.007**	Significant
CFA vs. Borderline-SMOTE	**0.002**	Significant

(p-value < 0.05 indicates statistical significance.)

6.5 Summary of Findings

Why CFA Outperforms Traditional Methods:

Aspect	SMOTE-Based Methods	CFA (Proposed Method)
Data Generation	Interpolates between samples	Uses real-world counterfactual pairs
Feature Preservation	May introduce unrealistic values	Retains meaningful feature importance (validated by SHAP)
Decision Boundary Learning	Can create overlapping class regions	Generates realistic, decision-boundary-aligned samples
False Negatives	Higher due to class overlap	Lower, improving recall and F1-score

7. Conclusion and Future Work

7.1 Conclusion

Class imbalance is a persistent challenge in machine learning, especially in critical applications such as **medical diagnosis, fraud detection, and cybersecurity**. Traditional oversampling techniques, such as **SMOTE, Borderline-SMOTE, and ADASYN**, attempt to mitigate this issue by generating synthetic minority samples. However, these methods often introduce **unrealistic interpolations, class overlap, and noise**, leading to poor generalization and increased false positives.

In this research, we introduced **Counterfactual Augmentation (CFA)** as a novel technique that leverages **observed counterfactual pairs to generate realistic synthetic minority instances**. Unlike SMOTE-based methods, CFA:

- **Uses real-world feature transformations** rather than artificial interpolations.
- **Ensures synthetic instances align with natural decision boundaries**, reducing class overlap.
- **Preserves meaningful feature importance**, validated through **SHAP-based analysis**.
- **Significantly improves minority class recall and F1-score**, with **5-10% gains over SMOTE and its variants**.

Extensive experiments on six benchmark datasets demonstrated that **CFA consistently outperforms traditional augmentation techniques**, showing statistically significant improvements in **recall, precision, and AUC-ROC scores**. These results indicate that CFA is an **effective and interpretable solution** for handling class imbalance in machine learning.

7.2 Limitations and Challenges

While CFA has demonstrated **strong performance improvements**, several limitations should be considered:

1. **Dependency on Counterfactual Pairs:**
    - CFA requires the identification of **naturally occurring counterfactual instances** in a dataset.
    - In some datasets, valid counterfactual pairs may be sparse or difficult to extract, requiring alternative pairing strategies.
2. **Computational Complexity:**
    - Unlike SMOTE, which applies simple interpolations, CFA involves **distance-based searches** and **feature transformations**, increasing computational overhead.
    - Optimizing CFA for **high-dimensional datasets** remains an area for future research.

3. **Generalization to Multi-Class Imbalance:**
    - This study primarily focused on **binary classification problems**.
    - Extending CFA to handle **multi-class imbalance scenarios** remains an open challenge.

Despite these limitations, CFA provides a **strong foundation for future advancements in counterfactual-based data augmentation**.

7.3 Future Work

7.3.1 Extending CFA to Deep Learning Models

- Current experiments were conducted on traditional classifiers (**Logistic Regression, Random Forest, k-NN, and MLP**).
- Future work will explore how **CFA can enhance deep learning architectures**, such as:
    - **CNNs (Convolutional Neural Networks)** for imbalanced image classification tasks.
    - **RNNs/LSTMs (Recurrent Neural Networks)** for time-series anomaly detection.

7.3.2 Multi-Class Counterfactual Augmentation

- CFA currently targets **binary class imbalance** problems.
- A possible extension is **multi-class CFA**, where counterfactual instances are derived across multiple class boundaries.
- Research will focus on **hierarchical and adaptive counterfactual selection** for **imbalanced multi-class learning**.

7.3.3 Hybrid Integration with Generative Models

- CFA can be **combined with generative models** such as **GANs (Generative Adversarial Networks)** to create **counterfactual-guided synthetic data**.
- This would allow for:
    - More flexible augmentation across **complex feature spaces**.
    - Increased robustness by **synthesizing hard-to-generate minority class instances**.

7.3.4 Real-World Deployment and Case Studies

- Applying CFA to **real-world production environments** will validate its impact beyond controlled experiments.
- Case studies will include:
    - **Healthcare AI systems** for rare disease diagnosis.
    - **Fraud detection pipelines** in financial transactions.
    - **Cybersecurity threat detection** using imbalanced attack datasets